

AI Digital Twin Platform.

Your AI drafts every response. A human decides what goes out. Nothing is ever sent, saved, or acted on without explicit approval.

Human-in-the-loop by design

Policy-constrained drafting

Full decision audit trail

LIVE TASK PREVIEW

Support Reply · AI Drafted

"Can I get a refund on order #12345?"

AI has drafted a policy-compliant response with a confidence score. Nothing sends until you decide.

Edit

Approve

• **HOW IT WORKS**

Three steps. Zero surprises.

Every inbound message follows the same simple, auditable path — from raw text to a fully human-approved outcome.

1

Submit

Paste an inbound message — a support ticket, a sales inquiry, a contract to review.

2

AI Drafts

A policy-constrained model produces a structured draft with a confidence score and risk flags — never a guess dressed up as certainty.

3

You Approve

Nothing goes out until a human reviews it. Approve, edit, reject, or escalate — you decide.

Authority, earned progressively

We never hand an AI more control than it has proven it deserves. Your team decides when — and if — to move to the next stage.

STAGE 1

Observation

AI recommends. It never acts.

STAGE 2 · CURRENT

Approval

AI drafts. A human decides.

STAGE 3

Delegation

AI executes within tight, pre-approved limits.

STAGE 4

Autonomous

Fully independent — reserved for narrow, proven workflows.

Built for real business workflows

Support Reply

Draft accurate, on-policy responses to customer support tickets in seconds.

Sales Email

Generate persuasive outreach and follow-ups without pricing or commitment risk.

Contract Review

Flag risky clauses and liability language before anything is signed.

• **CONTROL & OVERSIGHT**

You are always the last word.

Every AI draft resolves to exactly one of four explicit human decisions — nothing is silently auto-approved or auto-sent.

● **Approve**

The draft is sent exactly as written.

● **Edit & Approve**

You refine the draft, then send your version.

● **Reject**

The draft is discarded — nothing is sent.

● **Escalate**

Routed to a senior reviewer for judgment.

Measure what the AI is actually doing

A built-in eval dashboard tracks approval rate, edit rate, and quality over time — so trust in the system is based on evidence, not assumption.

Approval Rate

Share of AI drafts approved as-is, tracked over time and by workflow type.

Edit Rate

How often a draft needed a human correction before it was usable.

Quality Tags

Reviewer-rated Good / Needs Edit / Bad on every decision made.

Enforced guardrails — not a suggestion, a policy

Every AI draft is checked against a fixed, non-negotiable policy before it ever reaches a reviewer. Flagged examples include:

Max 15% discount without approval

"we guarantee"

"legally binding"

"100% guaranteed"

"unconditional refund"

"free forever"

"unlimited liability"

"auto-renew"



Full decision history

Every approval, edit, and rejection is logged and reviewable.



Nothing acts without you

No auto-send, no silent execution — ever, by design.



Evidence-based trust

Approval and edit rates make AI performance measurable, not assumed.

• GET STARTED

Ready to see it in action?

Create an account and submit your first task — the AI drafts, you decide. No commitment, no risk to your existing workflows.

1

Create your account

Sign up in seconds with your own work email — no setup call required.

2

Submit your first task

Paste a real support ticket, sales email, or contract to review.

3

Review the AI draft

Approve, edit, reject, or escalate — you're always in control.

GET STARTED AT

twin-platform-phi.vercel.app

CURRENT PLAN

Phase 1 — Approval Mode

No credit card required to create an account.